

# ECE472

## Deep Learning - Assignment 2

Yuval Epstain Ofek

September 16, 2020

### Summary

Spiral data was generated containing two spiral classes with 250 points each. The spirals were generated in polar form  $(r, \theta)$  by uniformly sampling across the range of  $\theta$ , the angle, and following the equation  $r = \theta - \theta_{initial} + r_{offset} + noise(SD = 0.3)$  to get the radius. The noise is random noise with a standard of deviation of 0.3,  $\theta_{initial}$  is the initial value of  $\theta$  for a given spiral, and the offset, set to 2, was used to separate the classes spatially from the origin. A multi-layer perceptron with 4 hidden layers of sizes [32, 16, 8, 8] was trained on this data to perform binary classification between the two spirals using the Adam optimizer. The loss during training was calculated using binary cross entropy with an  $L^2$  regularization penalty. After training, all training data was classified correctly. The training results are outlined in figure 1, which includes the spiral data and the learned boundary corresponding to  $p(t = 1|x) = 0.5$  as predicted by the perceptron.

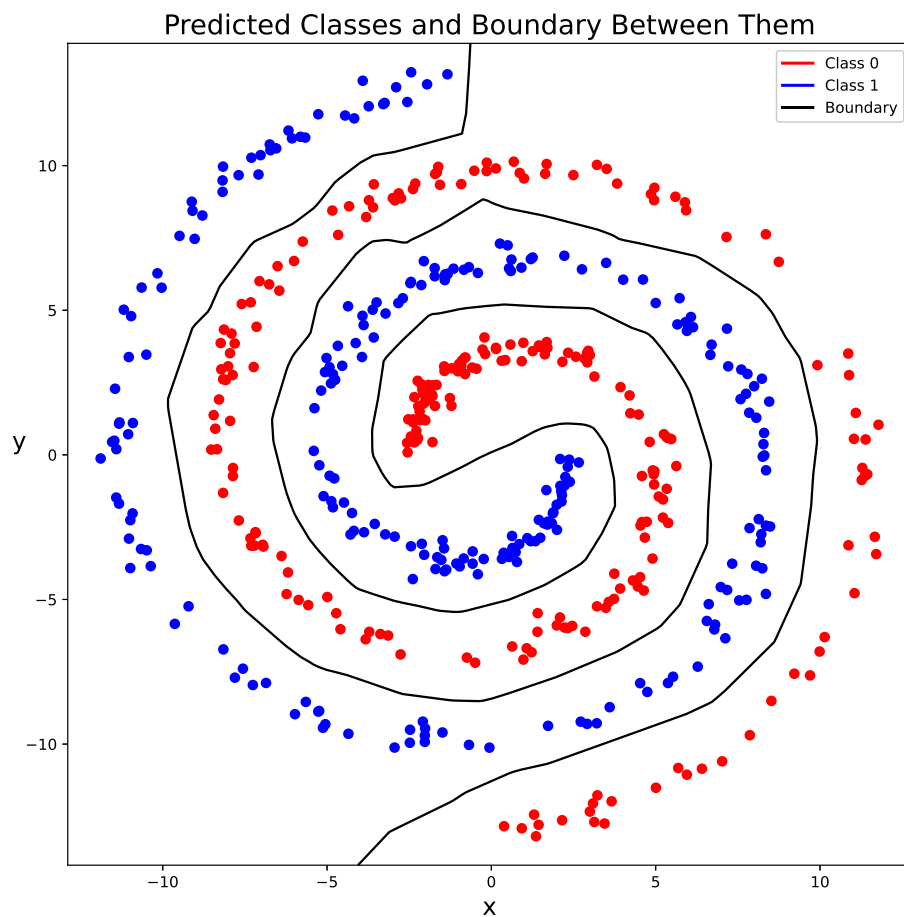


Figure 1: Perceptron Training Results on Spiral Dataset (500 points, random noise in radius with  $\sigma$  of 0.3). The boundary represents  $p(t = 1|x) = 0.5$  as predicted by the perceptron.

## My Various Attempts:

I went through a number of perceptron parameters before I got a set that worked with my data. When approaching the problem, the first thing I decided was that I will have relu activation for all the layers but the last one, which will have a sigmoid activation to match the binary output. This choice was made because I wanted to not have to worry about vanishing gradients. I also knew in the back of my head that I wanted to have a small number of weights.

My first attempt was a perceptron with one hidden layer of size 2048 (I generally felt that choosing powers of 2 was a good option). This didn't work at all, so since I felt like this was already a lot of "neurons", I decided to try out using multiple hidden layers, splitting the 2048 layer into two 1024 layers, followed by four layers of size 512 when that didn't work. At this point the results were very close to matching the dataset, so I decided to try cutting the number of "neurons" per layer, thinking that this may help restrict the perceptron whose boundary was a bit too over the place. I cut the hidden layers to size 64, which did the trick and at this point we had no errors.

I still felt like I could reduce the number of parameters even more, so I kept halving the layer sizes until I broke the model when we reached layers of size 16. I jumped back to the four hidden layers with size 32, and started decreasing the dimensions individually from the end of the perceptron (layer closest to the output) the input. The end result was a perceptron with 4 hidden layers of sizes [32, 16, 8, 8] in order from layers closest to the input to layers closest to the output.

## Appendix I- Python Code

```
# -*- coding: utf-8 -*-
"""DeepLearningAssignment2.ipynb

Automatically generated by Colaboratory.

Original file is located at
    https://colab.research.google.com/drive/1P38aS7q8eOUSWkoZSb6eHz9NNkZ-VwTv

**ECE472, Deep Learning - Assignment 2**

Submit by Sept. 16, 10pm

tldr: Perform binary classification on the spirals
dataset using a multi-layer perceptron.
You must generate the data yourself.

* I used the convention  $x$  in  $\mathbb{R}^2$  is the input vectors and  $y$ 
in  $\{0,1\}$  are labels
"""

import numpy as np
import tensorflow as tf
import numpy.random as npr
import matplotlib.pyplot as plt
import matplotlib

BATCH_SIZE = 64
#data parameters
N = 250
noise_sd = 0.3
offset = 2
seed = 4
theta_range1 = (np.pi/2, 4*np.pi)
theta_range2 = (3*np.pi/2, 5*np.pi)

#model parameters:
num_steps = 2_500
layer_sizes = [2,32, 16, 8, 8, 1]
lamdb = 0.01
colors = ['red','blue']

#Plotting Parameters:
```

```
N_bounds = 200
```

```
class spiral_data:
    def __init__(self, theta_range1, theta_range2,
                  noise_sd, offset, N1, shuffle = True,
                  N2 = None):
    def spiral(theta_range, noise_sd, offset, N):
        '''
        Generate N uniformly sampled points in a spiral with theta range
        theta_range with normal noise of variance noise_sd, and r offset from
        zero of offset.
        '''
        range_min, range_mx = theta_range
        theta = npr.uniform(range_min, range_mx, size = (N,1))
        r = offset + theta - range_min + npr.normal(size = (N,1))*noise_sd
        x_cord, y_cord = r*np.sin(theta), r*np.cos(theta)
        return np.concatenate((x_cord,y_cord), axis =1)

    #create 2 spirals given input parameters
    if N2 is None:
        N2 = N1
    spiral1 = spiral(theta_range1, noise_sd, offset, N1)
    spiral2 = spiral(theta_range2, noise_sd, offset, N2)

    self.x = np.concatenate((spiral1, spiral2), axis = 0)
    self.y = np.concatenate((np.ones(shape = (N1,1)),
                              np.zeros(shape = (N2,1))), axis = 0)

    if shuffle:
        shuffler = np.random.permutation(len(self.x))
        self.x = self.x[shuffler]
        self.y = self.y[shuffler]

    def get_batch(self, batch_size = BATCH_SIZE):
        '''
        return a batch of random batch_size features and their labels
        '''
        idxs = range(len(self.x))
        choices = npr.choice(idxs, batch_size)
        return (tf.convert_to_tensor(self.x[choices], dtype=tf.float32),
                tf.convert_to_tensor(self.y[choices], dtype=tf.float32))

class perceptron:
    def __init__(self,
                  layer_sizes,
```

```

        lambda,
        seed,
        initializer = tf.keras.initializers.GlorotNormal):
'''
initialize weights and model parameters
'''
#store parameters
self.layers = len(layer_sizes)
self.lambda = lambda

#initialize weights
initializer = initializer(seed)
params = {}
for i in range(1, self.layers):
    params['W'+str(i)] = tf.Variable(initializer(shape=(layer_sizes[i-1],
                                                         layer_sizes[i])))
    params['b'+str(i)] = tf.Variable(initializer(shape=(1,
                                                         layer_sizes[i])))
self.params = params

def layer_pass(self, x, W,b, activation):
'''
A pass through one layer with weights W and offset b, followed by
an activation function
'''
Z = tf.matmul(x,W)+b
if activation == "relu":
    A = tf.nn.relu(Z)
elif activation == 'sigmoid':
    A = tf.nn.sigmoid(Z)
else:
    A = Z
return A

def predict(self, x):
'''
predict y given input x and current model weights
'''
A = x
params = self.params
activation = 'relu'
for i in range(1, self.layers):
    A_prev = A
    if i+1 == self.layers:

```

```

        activation = 'sigmoid'
        A = self.layer_pass(A_prev, params['W'+str(i)],
                             params['b'+str(i)], activation)
    return A

def loss(self, y_hat, y):
    """
    binary cross entropy loss with L2 normalization
    """
    loss = tf.reduce_mean(-y*tf.math.log(y_hat)-(1-y)*tf.math.log(1-y_hat))
    #L2 regularization
    for i in range(1,self.layers):
        loss += tf.nn.l2_loss(self.params['W'+str(i)])*self.lambd
    return loss

def step(self,optimizer, cache):
    """
    A single step - predicting values, getting loss,
    & using the gradients of loss to update the
    parameters with Adam optimizer
    """
    with tf.GradientTape(persistent=True) as tape:
        x,y = d1.get_batch()
        tape.watch(self.params)
        y_hat = self.predict(x)
        lss = self.loss(y_hat,y)
    cache.append(lss)

    #get update paramaters
    grads = tape.gradient(lss, self.params)
    optimizer.apply_gradients(zip(grads.values(), self.params.values()))
    #return updated loss cache
    return cache

#run experiment
#generate data
npr.seed(seed)
d1 = spiral_data(theta_range1, theta_range2, noise_sd, offset, N)

#initialize perceptron
p = perceptron(layer_sizes,lambd,seed)

cache = []
optimizer = tf.optimizers.Adam()

```

```

for i in range(num_steps):
    cache = p.step(optimizer, cache)
    #Training printout
    if not((i+1)%100):
        p_train = int((i+1)/num_steps*100)
        y_hat = p.predict(tf.convert_to_tensor(d1.x, dtype=tf.float32))
        err = sum(abs(np.round(y_hat)-d1.y))[0]/len(d1.y)
        curr_loss = cache[-1].numpy()
        print(f'\r Training completed: {p_train}% |',
              f'Loss: {curr_loss} |',
              f'Test error: {err}',
              end = '')

#learning curve
plt.figure
plt.figure(figsize=(5,5))
plt.plot([lss.numpy() for lss in cache])
plt.title('Learning Curve', fontsize=14)
plt.xlabel('iteration', fontsize=12)
h = plt.ylabel('loss', fontsize=12)
h.set_rotation(0)
plt.show()

#final prediction plot
y_hat = p.predict(tf.convert_to_tensor(d1.x, dtype=tf.float32))

#boundary line
x_coord, y_coord = np.meshgrid(
    np.linspace(np.min(d1.x[:, 0])-1, np.max(d1.x[:, 0])+1, N_bounds),
    np.linspace(np.min(d1.x[:, 1])-1, np.max(d1.x[:, 1])+1, N_bounds),
    )
z = np.vstack((x_coord.flatten(), y_coord.flatten())).T
z_out = p.predict( tf.convert_to_tensor(z, dtype=tf.float32)).numpy()

#plotting
plt.figure
plt.figure(figsize=(10,10))
plt.scatter(d1.x[:,0], d1.x[:,1], c=np.round(y_hat.numpy(),0),
            cmap=matplotlib.colors.ListedColormap(colors))
plt.xlabel('x', fontsize=16)
h = plt.ylabel('y', fontsize=16)
h.set_rotation(0)
plt.title('Predicted Classes and Boundary Between Them', fontsize=18)
plt.contour(x_coord, y_coord,

```



```

        np.squeeze(np.reshape(z_out, x_coord.shape)),
        levels=1, colors="k", )
plt.savefig('Prediction.eps', format='eps')
from matplotlib.lines import Line2D
cmap = plt.cm.coolwarm
custom_lines = [Line2D([0], [0], color='r', lw=2),
                 Line2D([0], [0], color='b', lw=2),
                 Line2D([0], [0], color='k', lw=2)]

plt.legend(custom_lines, ['Class 0', 'Class 1', 'Boundary'])
plt.show()

```